

# Reliability of Explainable AI (XAI)

October 19, 2025

## 1 Abstract

**Problem Statement:** Quantifying the *reliability* of feature-attribution Explainable AI (XAI) for **software defect prediction**—i.e., whether explanations truly reflect how the trained model behaves.

**Proposed Methodology:** We introduce a general evaluation framework that compares explainer rankings to a *grounded behavioral signal* derived from the model itself. Concretely, we (i) estimate *true loss sensitivity* via small  $\varepsilon$ -perturbations and finite differences on binary cross-entropy, (ii) assess **Gradient–Loss Reliability (GLR)** using Spearman rank correlation between SHAP importances and normalized sensitivity, (iii) validate *action-consistency* with  $\varepsilon$ -**Hit@k**, and (iv) quantify probability honesty with **Expected Calibration Error (ECE)**. A composite **ReliabilityScore** summarizes these facets.

**Application:** We demonstrate the framework on the **NASA PROMISE CM1** dataset (module-level defect prediction) using a **Random Forest** classifier with **SHAP (TreeExplainer)**. The pipeline reports discrimination (AUC, F1, Precision, Recall, Brier), GLR,  $\varepsilon$ -Hit@k, ECE, and the composite ReliabilityScore, alongside visualization (top mean $|\phi|$  features, GLR histograms, calibration curves).

**Note:** Software defect prediction aims to identify defect-prone modules from static code metrics (size, complexity, Halstead measures). Reliable explanations help practitioners prioritize inspections and reason about model decisions in safety- and cost-sensitive engineering workflows.

## 2 Introduction

**Problem.** Machine Learning (ML) is widely used for **software defect prediction** from static code metrics. However, many high-performing models (e.g., Random Forests, Neural Networks) behave as *black boxes*: they output probabilities without revealing *why*. In software engineering workflows—where inspection budgets, safety, and maintenance decisions depend on trust—engineers need transparent, dependable reasoning behind predictions.

**Role of XAI.** Explainable AI (XAI) provides per-feature attributions that indicate which metrics (e.g., LOC, cyclomatic complexity, Halstead measures) drive a module’s defect risk. In this work we use **SHAP (TreeExplainer)** to obtain local and global importance over code metrics.

**Issue with Current Explanations.** Explanations can be *unstable* across folds, data slices, or minor distribution shifts, and they may not faithfully reflect the model’s *true behavior* near a prediction. Without a standard reliability check, feature attributions risk being misleading or overconfident.

**Contribution.** We propose a **structured reliability evaluation** for feature-attribution XAI on defect prediction:

- A *behavior-grounded signal* via small  $\varepsilon$ -perturbations and binary cross-entropy to estimate **true loss sensitivity**.

- **Gradient–Loss Reliability (GLR)**: Spearman rank alignment between SHAP importances and normalized true-loss sensitivities.
- **$\varepsilon$ -Hit@k**: an action-consistency test that checks whether SHAP’s top- $k$  features also cause the largest loss increases under tiny nudges.
- **Calibration (ECE)**: a probability honesty check complementary to discrimination metrics.
- A composite **ReliabilityScore** aggregating GLR,  $\varepsilon$ -Hit@k, and 1–ECE for a single, interpretable summary.

We demonstrate this framework on the **NASA PROMISE CM1** dataset with a **Random Forest** model and SHAP explanations, reporting discrimination (AUC, F1, Precision, Recall, Brier), explanation reliability (GLR,  $\varepsilon$ -Hit@k), and calibration (ECE), alongside diagnostic visualizations.

### 3 Literature Survey

- **Emphasis gap.** Most prior work prioritizes *accuracy* or raw *interpretability*, while *reliability* is rarely stress-tested. Existing reliability implementations are largely *domain-specific*, often framed as a single generalized  $k$ -fold SHAP/LIME pipeline.
- **Under-specified dimensions.** Traditional approaches under-specify core explainability dimensions: explanations *vary with train/test splits* (hurting generalization), can *disagree with model attributions* (low concordance), and may *lack stability*.

Title of the Paper	Published In	Significance	Limitation
<b>LCNN: Lightweight CNN for Software Defect Feature Identification Using Explainable AI</b>	IEEE Access 2024	Used SHAP and LIME to interpret CNN-based predictions, highlighting reproducibility of explanations in deep models.	No systematic reliability assessment (e.g., stability, concordance); limited to visualization of explanations.
<b>A Methodology for Reliability Analysis of Explainable AI: Application to Endocrinology Diseases</b>	IEEE Access 2024	Introduced a generalized framework with <i>Generalizability</i> , <i>Concordance</i> , and <i>Stability</i> as reliability metrics for SHAP/LIME.	Domain-specific evaluation (healthcare); ignores local sensitivity / perturbation-based reliability.
<b>Understanding Software Defect Prediction Through eXplainable Neural Additive Models (XNAMs)</b>	IEEE Access 2025	Designed inherently interpretable models, reducing reliance on post-hoc explanations for consistency.	Focused on interpretability, not reliability; no cross-fold stability or perturbation analysis.

Table 1: Snapshot of recent work: strengths and reliability gaps.

Across these studies, there is still no universal standard for evaluating *explanation reliability*; cross-fold stability, concordance with model behavior, and perturbation-grounded checks remain under-explored and inconsistently reported.

## Current Scenario

### 4 Methodology

We propose to combine XAI with k-fold cross-validation and evaluate using three metrics: **Generalizability**, **Concordance**, and **Stability**. These are aggregated into a **Global Reliability Score**.

#### 4.1 Step 1: XAI with k-Fold Cross Validation

1. Split dataset into  $k$  folds.
2. For each fold:
  - Train the model on  $k - 1$  folds.
  - Test on the remaining fold.
  - Generate feature importance via SHAP or LIME.
3. Collect  $k$  sets of feature importances across folds.

This captures the variability of feature importance across different splits.

#### 4.2 Step 2: Define Key Metrics

**Generalizability.** Do train vs test explanations rank features similarly (per fold)?

Let  $\mathbf{r}_{\text{train}}^{(j)}$  and  $\mathbf{r}_{\text{test}}^{(j)}$  be the rank vectors (of mean  $|\text{SHAP}|$  importances) for fold  $j$  over the same  $n$  features. Compute Spearman's  $\rho$  and rescale to  $[0, 1]$ :

$$\rho_{\text{gen}}^{(j)} = \text{Spearman}(\mathbf{r}_{\text{train}}^{(j)}, \mathbf{r}_{\text{test}}^{(j)}), \quad G = \frac{1}{k} \sum_{j=1}^k \tilde{\rho}_{\text{gen}}^{(j)}, \quad \tilde{\rho} = \frac{\rho + 1}{2} \in [0, 1].$$

**Concordance.** Does the XAI importance agree with the model's internal (Gini/MDI) importance?

Let  $\mathbf{r}_{\text{SHAP}}^{(j)}$  and  $\mathbf{r}_{\text{Gini}}^{(j)}$  be rank vectors for fold  $j$ . Use **Spearman** (not Pearson), rescaled to  $[0, 1]$ :

$$\rho_{\text{con}}^{(j)} = \text{Spearman}(\mathbf{r}_{\text{SHAP}}^{(j)}, \mathbf{r}_{\text{Gini}}^{(j)}), \quad C = \frac{1}{k} \sum_{j=1}^k \tilde{\rho}_{\text{con}}^{(j)}, \quad \tilde{\rho} = \frac{\rho + 1}{2}.$$

**Stability.** Are explanations consistent across folds?

Let  $\{\mathbf{s}^{(j)}\}_{j=1}^k$  be the per-fold importance vectors (e.g., mean  $|\text{SHAP}|$  on *train*). Compute mean pairwise **Spearman** across all fold pairs, rescaled to  $[0, 1]$ :

$$S = \frac{2}{k(k-1)} \sum_{1 \leq p < q \leq k} \tilde{\rho}(\text{rank}(\mathbf{s}^{(p)}), \text{rank}(\mathbf{s}^{(q)})), \quad \tilde{\rho} = \frac{\rho + 1}{2}.$$

#### 4.3 Step 3: Global Reliability Score

We combine the three metrics via their **arithmetic mean**:

$$\text{Reliability} = R = \frac{G + C + S}{3},$$

with  $R \in [0, 1]$ ; larger values indicate more reliable explanations.

## 5 Implementation

### 5.1 Generalizability

Generalizability measures how well the feature importance derived from the training data generalizes to unseen test data. It is quantified using the **Spearman rank correlation coefficient** between the ranks of feature importances obtained from the training and test sets in each fold of cross-validation.

**Formula.** Let  $R_{train}^{(j)}$  and  $R_{test}^{(j)}$  be the rank vectors of feature importances (e.g., mean absolute SHAP values) for fold  $j$ . The Spearman correlation coefficient  $\rho$  for fold  $j$  is given by:

$$\rho^{(j)} = 1 - \frac{6 \sum_{i=1}^n d_i^2}{n(n^2 - 1)}$$

where

- $n$  is the number of features,
- $d_i = \text{rank}_{train}(f_i) - \text{rank}_{test}(f_i)$  is the difference between the training and testing ranks of feature  $f_i$ .

The overall generalizability score across  $k$  folds is the average:

$$G = \frac{1}{k} \sum_{j=1}^k \rho^{(j)}$$

**Example.** Consider  $n = 4$  features with the following ranks:

Feature	Train Rank	Test Rank	$d_i$	$d_i^2$
$f_1$	1	1	0	0
$f_2$	2	3	-1	1
$f_3$	3	2	1	1
$f_4$	4	4	0	0

Here,  $\sum d_i^2 = 2$ . Substituting in the formula:

$$\rho = 1 - \frac{6 \cdot 2}{4 \cdot (4^2 - 1)} = 1 - \frac{12}{60} = 0.8$$

Thus, the generalizability for this fold is 0.8, indicating a high consistency between train and test feature rankings.

### 5.2 Concordance

Concordance measures the agreement between feature importance obtained from post-hoc explainability methods and those derived directly from the learning algorithm. In our case, we compare **SHAP-based feature importances** with the **Gini importance** (Mean Decrease in Impurity, MDI) computed from tree-based models such as Random Forest. The agreement is quantified using the **Spearman rank correlation coefficient**.

**Formula.** Let  $R_{\text{SHAP}}^{(j)}$  and  $R_{\text{Gini}}^{(j)}$  be the rank vectors of feature importances obtained by SHAP and Gini index, respectively, for fold  $j$ . The Concordance for fold  $j$  is:

$$C^{(j)} = \rho(R_{\text{SHAP}}^{(j)}, R_{\text{Gini}}^{(j)})$$

where  $\rho$  is the Spearman correlation coefficient defined as:

$$\rho = 1 - \frac{6 \sum_{i=1}^n d_i^2}{n(n^2 - 1)}$$

with

- $n$  being the total number of features,
- $d_i = \text{rank}_{\text{SHAP}}(f_i) - \text{rank}_{\text{Gini}}(f_i)$  being the difference in rank of feature  $f_i$  under SHAP and Gini importance.

The overall concordance score across  $k$  folds is then:

$$C = \frac{1}{k} \sum_{j=1}^k C^{(j)}$$

**Example.** Suppose we have four features with the following importance ranks:

Feature	SHAP Rank	Gini Rank	$d_i$	$d_i^2$
$f_1$	1	2	-1	1
$f_2$	2	1	1	1
$f_3$	3	3	0	0
$f_4$	4	4	0	0

Here,  $\sum d_i^2 = 2$  and  $n = 4$ . Substituting:

$$\rho = 1 - \frac{6 \cdot 2}{4 \cdot (4^2 - 1)} = 1 - \frac{12}{60} = 0.8$$

Thus, the concordance for this fold is 0.8, indicating strong alignment between SHAP and Gini-based explanations.

### 5.3 Stability

Stability quantifies how consistent the *explanations* (feature-importance vectors) are across different resamples/folds of the data. Given one importance vector per fold (e.g., mean absolute SHAP values on the *train* split of each fold), we compute the **Spearman rank correlation** for every pair of folds and average these correlations.

**Setup.** Let  $F$  be the number of folds and let  $\mathbf{s}^{(p)} \in \mathbb{R}^n$  denote the vector of feature importances for fold  $p$  over  $n$  common features. Let  $\mathbf{r}^{(p)} = \text{rank}(\mathbf{s}^{(p)})$  be the rank vector (1 = most important). Then the Spearman correlation between folds  $p$  and  $q$  is

$$\rho(\mathbf{r}^{(p)}, \mathbf{r}^{(q)}) = 1 - \frac{6 \sum_{i=1}^n d_i^2}{n(n^2 - 1)}, \quad d_i = r_i^{(p)} - r_i^{(q)}.$$

**Definition.** The Stability score is the mean pairwise Spearman correlation:

$$S = \frac{2}{F(F-1)} \sum_{1 \leq p < q \leq F} \rho(\mathbf{r}^{(p)}, \mathbf{r}^{(q)}).$$

**Example.** Assume  $F = 3$  folds and  $n = 4$  features. The Spearman correlations between fold pairs are  $\rho_{12} = 0.86$ ,  $\rho_{13} = 0.78$ ,  $\rho_{23} = 0.80$ . Then

$$S = \frac{2}{3 \cdot 2} (0.86 + 0.78 + 0.80) = \frac{1}{3} \cdot 2.44 \approx 0.813.$$

## 5.4 Reliability

Reliability provides an overall measure of trustworthiness by aggregating the three criteria: Generalizability ( $G$ ), Concordance ( $C$ ), and Stability ( $S$ ). Since all three values lie in the range  $[0, 1]$ , we define the reliability score as their arithmetic mean:

$$R = \frac{G + C + S}{3}$$

where

- $G$  is the generalizability score,
- $C$  is the concordance score,
- $S$  is the stability score.

**Example.** Suppose we obtained:

$$G = 0.75, \quad C = 0.80, \quad S = 0.60$$

Then:

$$R = \frac{0.75 + 0.80 + 0.60}{3} = \frac{2.15}{3} \approx 0.717$$

Thus, the overall reliability of the explanations is 0.717, indicating moderately strong reliability.

## 6 Results

Table 2: Reliability metrics (scaled to  $[0, 1]$ ).

Metric	Value
Generalizability (Spearman train vs. test)	0.9866
Concordance (SHAP vs. MDI, Spearman)	0.9154
Stability (Pairwise Spearman across folds)	0.7788
Reliability Index ( $\frac{G+C+S}{3}$ )	0.9468

### Insights.

- **Generalizability (0.9866):** Train–test explanation ranks are almost identical, indicating the explanation pattern holds strongly on unseen data.
- **Concordance (0.9154):** High agreement between SHAP and model-internal Gini (MDI) importances; post-hoc XAI is consistent with the model’s own attributions.
- **Stability (0.7788):** Explanation vectors are fairly consistent across folds (mean pairwise Spearman), with some expected variability from resampling.
- **Reliability Index (0.9468):** Averaging  $G$ ,  $C$ , and  $S$  yields a strong overall reliability—explanations are robust, model-aligned, and reasonably stable.

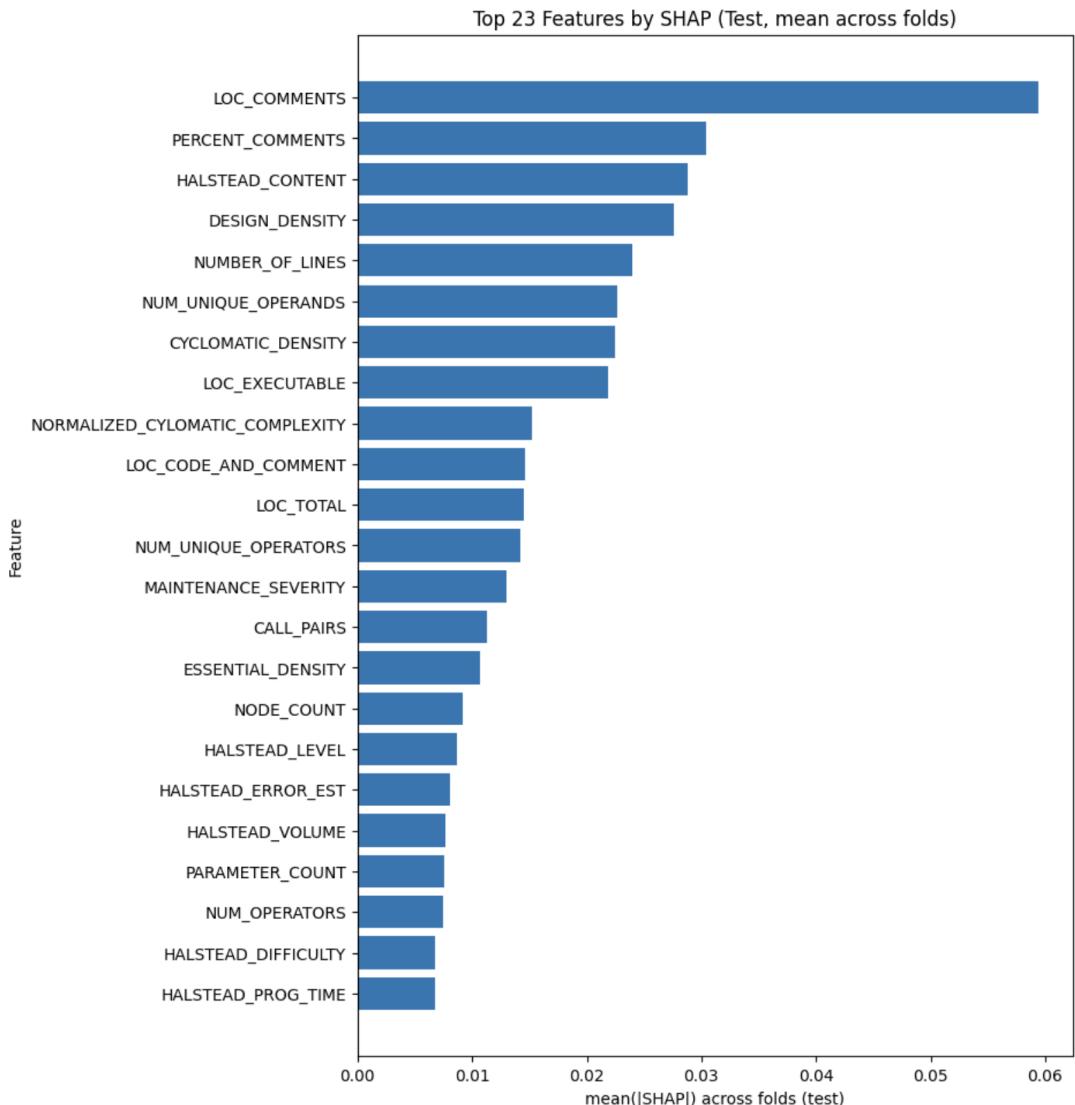


Figure 1: Visualization of feature importance scores using SHAP (SHapley Additive exPlanations) for the software defect dataset. Each bar corresponds to the average absolute SHAP value of a feature, representing its global contribution to the model’s predictive performance. Features with longer bars have a greater influence on the model’s decision-making process. This visualization highlights which software metrics (e.g., lines of code, number of developers, modification entropy) play the most critical role in predicting defect proneness.

**Explanation.** The SHAP method is grounded in cooperative game theory, where each feature is considered a “player” contributing to the final prediction. The SHAP value for a feature quantifies its marginal contribution averaged across all possible feature coalitions. By plotting the mean absolute SHAP values, we obtain a ranking of features by importance. This allows us to interpret the model globally, identifying the most influential factors driving defect prediction. For instance, in many defect datasets, features like *code churn*, *developer experience*, and *complexity metrics* emerge as dominant contributors, aligning with empirical software engineering studies.

## Proposed Method

### 7 Framework

The proposed framework evaluates the **reliability of XAI** for a **software defect prediction** model trained with a Random Forest, using SHAP explanations and three reliability checks: *GLR* (*rank alignment*),  $\varepsilon$ -*Hit@k* (perturbation validity), and *ECE* (probability calibration).

#### 7.1 Dataset

- Software defect dataset with binary target `Defective` (e.g., `CM1.csv`).

#### 7.2 Data Preprocessing (as in code)

- Load CSV; split into features  $X$  and label  $y = \text{Defective}$ .
- Drop constant columns:** remove any feature with only one unique value.
- Missing values:** if present, apply forward fill then back fill.
- Note:* No scaling/standardization is applied in the code.

#### 7.3 Model Training & Evaluation

- Cross-validation:** StratifiedKFold with  $n\_splits$  folds (shuffle, fixed seed).
- Class imbalance (train only):** SMOTE with adaptive  $k$ -neighbors:

$$k = \min(5, \max(1, \#\text{minority} - 1)).$$

- Classifier:** RandomForestClassifier

```
n_estimators = 600, min_samples_leaf = 2, class_weight = "balanced_subsample", random_stan
```

- Per-fold metrics on test:** AUC, F1, Precision, Recall, Brier score.
- Aggregate means across folds:  $\overline{\text{AUC}}$ ,  $\overline{\text{F1}}$ ,  $\overline{\text{Precision}}$ ,  $\overline{\text{Recall}}$ ,  $\overline{\text{Brier}}$ .

#### 7.4 Explainable AI (XAI) with SHAP (as in code)

- Use `TreeExplainer` with a background sample of  $\min(256, |X_{\text{train}}|)$ , `feature_perturbation=interventional`, `model_output=probability`.
- Compute per-sample SHAP values on test; use absolute values  $|\phi|$ .
- Compute global importance as mean of  $|\phi|$  across test samples.

#### 7.5 True Loss Sensitivity via Finite Differences (as in code)

- Binary cross-entropy on probability  $p$ :

$$\ell(p, y) = -\left(y \log p + (1 - y) \log(1 - p)\right).$$

- (b) For each test sample  $x \in \mathbb{R}^d$  and each feature  $j$ , approximate

$$g_j \approx \left| \frac{\partial \ell}{\partial x_j} \right| \approx \frac{|\ell(p(x + \varepsilon e_j), y) - \ell(p(x - \varepsilon e_j), y)|}{2\varepsilon},$$

where  $p(\cdot)$  is `predict_proba` class-1 output of the RF and  $e_j$  is the unit vector.

- (c) Normalize per sample:

$$\tilde{g}_j = \frac{g_j}{\sum_{k=1}^d g_k + \epsilon}.$$

## 7.6 Gradient–Loss Reliability (GLR)

- (A) For each test sample, take vectors  $S$  (absolute SHAP) and  $G$  (normalized sensitivities  $\tilde{g}$ ).
- (B) Convert to ranks and compute **Spearman** correlation  $\rho$  between  $S$  and  $G$ :

$$\rho = 1 - \frac{6 \sum_{j=1}^d (r_j^S - r_j^G)^2}{d(d^2 - 1)}.$$

- (C) **GLR score** is the mean of  $\rho$  over all test samples (ignoring NaN).

## 7.7 $\varepsilon$ -Hit@k (Perturbation Validity Check, as in code)

- (A) For each selected test sample (random  $\sim 15\%$  per fold until a global budget `max_eps_samples`):
- Let Top- $M_{\text{SHAP}}$  be indices of the  $M = \min(\text{m\_limit}, d)$  largest  $|\phi|$ .
  - Let  $T_k$  be SHAP top- $k$  within those  $M$ .
  - For each  $j \in \text{Top-}M_{\text{SHAP}}$ , perturb  $x_j \leftarrow x_j + \varepsilon$ ; compute

$$\Delta\ell_j = \ell(p(x + \varepsilon e_j), y) - \ell(p(x), y).$$

- (d) Let  $D_k$  be indices of the top- $k$  features by  $\Delta\ell_j$ .

- (B) Define per-sample hit:

$$\text{Hit}@k(\text{sample}) = \mathbf{1}\{T_k \cap D_k \neq \emptyset\}.$$

- (C)  $\varepsilon$ -**Hit@k** is the mean of these hits over all sampled test rows.

## 7.8 Calibration (ECE, as in code)

- (a) Expected Calibration Error with  $B$  bins:

$$\text{ECE} = \sum_{b=1}^B \frac{n_b}{N} |\text{acc}(b) - \text{conf}(b)|,$$

where  $\text{conf}(b)$  is the mean predicted probability and  $\text{acc}(b)$  is the empirical positive rate in bin  $b$ .

## 7.9 Composite Reliability Score (as in code)

- Rescale GLR to  $[0, 1]$ :  $\rho' = \frac{\rho+1}{2}$ .

- Combine:

$$\text{ReliabilityScore} = \text{mean}\left(\rho', \text{Hit}@k, \max\{0, 1 - \text{ECE}\}\right).$$

## 7.10 Visualization (as produced by code)

- **Top features by mean(—SHAP—) across folds (horizontal bar plot).**
- **GLR histogram** showing distribution of per-sample Spearman  $\rho$ .
- **Calibration curve** (predicted probability vs. fraction of positives) with ECE in title.
- (Optional) **SHAP Beeswarm snapshot** for the last fold.

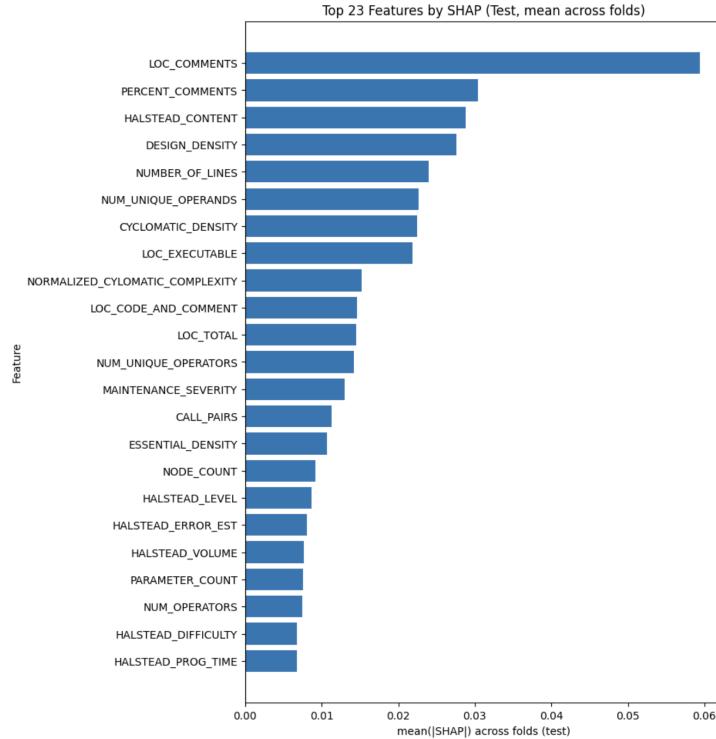


Figure 2: Top features by mean absolute SHAP on test (averaged across folds).

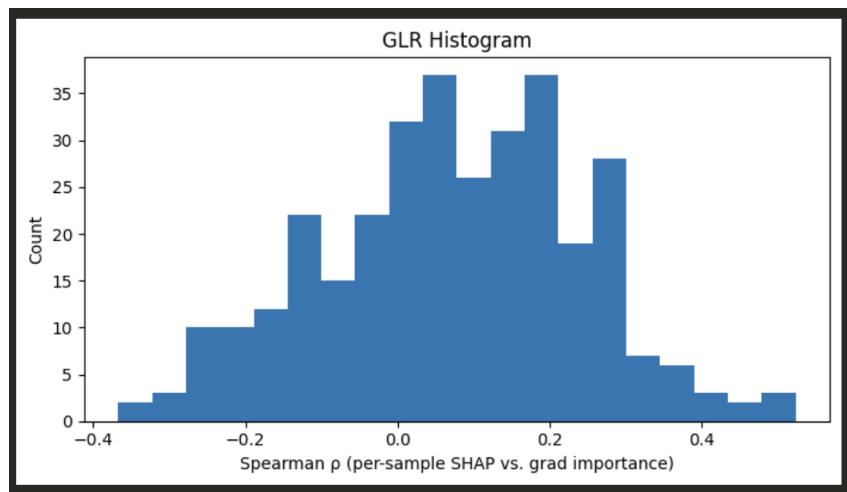


Figure 3: GLR histogram: per-sample Spearman  $\rho$  between SHAP and normalized sensitivity.

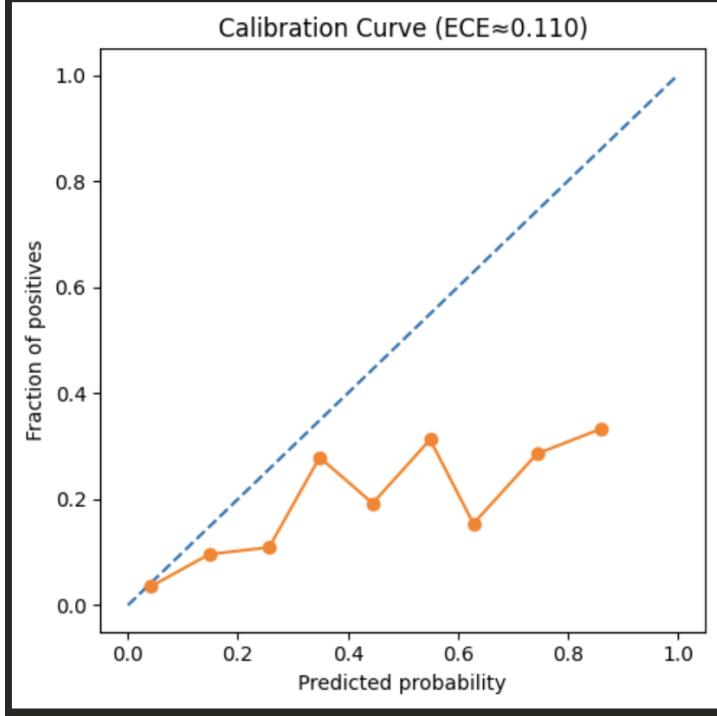


Figure 4: Calibration curve with ECE annotation.

### 7.11 Hyperparameters & Defaults (as in code)

- $\varepsilon = 10^{-3}$ , `top_k` = 10, `m_limit` = 20, `calib_bins` = 10.
- `max_eps_samples` = 50 (global budget for  $\varepsilon$ -Hit@k sampling).
- `n_splits` = 5, `rng` = 42.

## 8 Dataset and Problem Setting

We use the **NASA PROMISE CM1** software defect dataset (`cm1.csv`). CM1 corresponds to a NASA spacecraft instrument (C code) and is widely studied for **module-level defect prediction**. Each row is a software module described by static code metrics; the binary target is **Defective** (1 = defective, 0 = non-defective).

### Size and label.

- **Instances (modules):** 327
- **Attributes (features):** 35 numeric metrics (listed below)
- **Target (class):** Defective (1 or 0); in our file: 1 = 42, 0 = 285

### Attribute list (35 features).

1. `BRANCH_COUNT` — number of branches in control flow.
2. `CALL_PAIRS` — number of function call pairs.
3. `LOC_CODE_AND_COMMENT` — lines containing both code and comments.
4. `LOC_COMMENTS` — comment-only lines.

5. CONDITION\_COUNT — count of boolean conditions.
6. CYCLOMATIC\_COMPLEXITY — McCabe cyclomatic complexity.
7. CYCLOMATIC\_DENSITY — normalized cyclomatic complexity (complexity density).
8. DECISION\_COUNT — number of decision points (e.g., `if`, `case`).
9. DESIGN\_COMPLEXITY — design-level McCabe complexity.
10. DESIGN\_DENSITY — normalized design complexity (density).
11. EDGE\_COUNT — number of edges in the control-flow graph.
12. ESSENTIAL\_COMPLEXITY — essential (structuredness) complexity.
13. ESSENTIAL\_DENSITY — normalized essential complexity (density).
14. LOC\_EXECUTABLE — executable (code) lines.
15. PARAMETER\_COUNT — number of function/method parameters.
16. HALSTEAD\_CONTENT — Halstead content metric.
17. HALSTEAD\_DIFFICULTY — Halstead difficulty.
18. HALSTEAD EFFORT — Halstead effort.
19. HALSTEAD\_ERROR\_EST — Halstead estimated errors (bugs).
20. HALSTEAD\_LENGTH — Halstead length.
21. HALSTEAD\_LEVEL — Halstead level (inverse difficulty).
22. HALSTEAD\_PROG\_TIME — Halstead programming time estimate.
23. HALSTEAD\_VOLUME — Halstead volume.
24. MAINTENANCE\_SEVERITY — maintenance severity indicator.
25. MODIFIED\_CONDITION\_COUNT — count of modified conditions (MC/DC-related).
26. MULTIPLE\_CONDITION\_COUNT — count of multi-part conditions.
27. NODE\_COUNT — number of nodes in the control-flow graph.
28. NORMALIZED\_CYCLOMATIC\_COMPLEXITY — scaled cyclomatic complexity.
29. NUM\_OPERANDS — total operands (Halstead).
30. NUM\_OPERATORS — total operators (Halstead).
31. NUM\_UNIQUE\_OPERANDS — unique operands (Halstead).
32. NUM\_UNIQUE\_OPERATORS — unique operators (Halstead).
33. NUMBER\_OF\_LINES — total number of lines.
34. PERCENT\_COMMENTS — percentage of comment lines.
35. LOC\_TOTAL — total lines of code (LOC).

These metrics jointly capture size/volume (e.g., LOC, Halstead length/volume), structural and control-flow properties (e.g., nodes/edges/branches), and different notions of complexity (cyclomatic, essential, design), all of which are known to correlate with defect proneness.

[NASA PROMISE CM1 Dataset \(cm1.csv\)](#)

## 9 Model and Training Controls

**Model.** We use a **Random Forest** classifier (tree ensemble) rather than an MLP. Key hyperparameters (as used in code):

```
n_estimators=600,
min_samples_leaf=2,
max_depth=None,
class_weight="balanced_subsample",
n_jobs=-1.
```

Each fold uses a fixed `random_state` (fold id).

**Training control.** Random Forests do not use learning-rate schedules or early stopping. Regularization is controlled by ensemble size and tree shape (`min_samples_leaf`, `max_depth`). We report standard generalization metrics per fold (AUC, F1, Precision, Recall, Brier).

**Class imbalance.** We address imbalance via two mechanisms aligned with the code:

- (i) **SMOTE** oversampling is applied on *training folds only* to synthesize minority-class samples;
- (ii) **class weighting** with `balanced_subsample`, which re-weights classes within each bootstrap sample. (No `pos_weight` or logits-based BCE are used for training, since RF is non-parametric.)

**Explainability setup.** We compute per-sample SHAP attributions on test data using **TreeExplainer** with an interventional background of size  $\min(256, |X_{\text{train}}|)$  and `model_output = probability`. Global importance is the mean absolute SHAP over test samples.

## BCE on Predicted Probability (used in reliability checks)

Although the Random Forest is not trained with BCE, we use the **binary cross-entropy on probabilities**  $\ell(p, y)$  as a smooth, model-agnostic loss to probe reliability:

$$\ell(p, y) = -[y \log p + (1 - y) \log(1 - p)], \quad p = \Pr(y = 1 \mid x).$$

In code,  $p$  is clipped to  $[10^{-8}, 1 - 10^{-8}]$  for numerical stability.

## True-Loss Sensitivity via $\varepsilon$ -Perturbations

For a sample  $x \in \mathbb{R}^d$  and feature  $j$ , we estimate the magnitude of the loss sensitivity by a central finite difference:

$$g_j \approx \left| \frac{\partial \ell}{\partial x_j} \right| \approx \frac{|\ell(p(x + \varepsilon e_j), y) - \ell(p(x - \varepsilon e_j), y)|}{2\varepsilon},$$

with small  $\varepsilon = 10^{-3}$ . We then normalize per sample,

$$\tilde{g}_j = \frac{g_j}{\sum_{k=1}^d g_k + \epsilon},$$

to obtain a comparable importance profile used to (i) compute **GLR** (Spearman rank alignment between SHAP and  $\tilde{g}$ ) and (ii) validate explanations via  **$\varepsilon$ -Hit@k** (do SHAP top- $k$  features also yield the largest  $\Delta\ell$  under  $+\varepsilon$  nudges?).

## Calibration and Reliability Aggregation

Calibration quality is summarized by the **Expected Calibration Error** (ECE) computed over uniform probability bins. We also report a composite **ReliabilityScore** that averages a rescaled GLR,  $\varepsilon$ -Hit@k, and  $(1 - \text{ECE})$  (clipped at 0), providing a single, interpretable summary of explanation alignment, perturbation validity, and probability honesty.

## 10 Evaluation: Discrimination, Thresholding, Calibration

**Discrimination (AUC).** We compute the **ROC AUC** on test-fold probabilities from the Random Forest and report the *mean across folds*. AUC is threshold-free and summarizes ranking quality of positives vs. negatives.

**Thresholding (0.5) & Class Metrics.** Hard predictions are obtained via a fixed threshold  $\hat{y} = \mathbf{1}\{p \geq 0.5\}$ . We report **F1**, **Precision**, and **Recall** per fold on the test split, then average across folds. (Accuracy and Youden's  $J$  are *not* used in this pipeline.)

**Calibration: Brier score & ECE.** We assess probability quality with the **Brier score** and the **Expected Calibration Error** (ECE), using uniform probability bins. Let  $p_i$  be the predicted probability for sample  $i$  and  $y_i \in \{0, 1\}$  the true label.

$$\text{Brier} = \frac{1}{N} \sum_{i=1}^N (p_i - y_i)^2.$$

Let the predictions be partitioned into  $B$  bins; in bin  $b$ ,  $\text{conf}(b)$  is the mean predicted probability and  $\text{acc}(b)$  the empirical positive rate. Then

$$\text{ECE} = \sum_{b=1}^B \frac{n_b}{N} |\text{acc}(b) - \text{conf}(b)|.$$

We also form *reliability diagrams* via `calibration_curve` (uniform bins), plotting fraction of positives vs. mean predicted probability, and annotate the curve with the computed ECE.

**Fold Aggregation.** Per-fold metrics (AUC, F1, Precision, Recall, Brier) are computed on each test fold and averaged. For calibration plots and ECE, all test-set probabilities and labels from all folds are concatenated before binning to obtain a single overall reliability curve.

## 11 Explanation Methods

### 11.1 SHAP with Interventional TreeExplainer (as used in code)

We compute SHAP values on the test split with **TreeExplainer** (Random Forest):

- **Background:** a random training-fold sample of size  $\min(256, |X_{\text{train}}|)$  (no K-means), with `feature_perturbation=interventional`.
- **Model output:** `probability`. For binary tasks, we take the positive-class SHAPs and use  $|\phi|$ .
- **Global importance:** mean of  $|\phi|$  across test samples; we visualize top features by  $\text{mean}|\phi|$ .

## 11.2 True-Loss Sensitivity via $\varepsilon$ -Perturbations (Finite Differences)

**Setup.** Let  $x \in \mathbb{R}^d$  be a test sample,  $y \in \{0, 1\}$  its label, and  $p(x) = \Pr(y = 1 \mid x)$  the Random Forest probability. We use the **binary cross-entropy** on probability

$$\ell(p, y) = -(y \log p + (1 - y) \log(1 - p)),$$

with  $p$  clipped to  $[10^{-8}, 1 - 10^{-8}]$  for stability (as in code).

**Finite-difference sensitivity per feature  $j$ .** For a small  $\varepsilon > 0$  (code uses  $\varepsilon = 10^{-3}$ ),

$$g_j \approx \left| \frac{\partial \ell}{\partial x_j} \right| \approx \frac{|\ell(p(x + \varepsilon e_j), y) - \ell(p(x - \varepsilon e_j), y)|}{2\varepsilon},$$

where  $e_j$  is the  $j$ -th unit vector.

**Per-sample normalization.**

$$\tilde{g}_j = \frac{g_j}{\sum_{k=1}^d g_k + \epsilon},$$

yielding a distribution  $\tilde{g}$  over feature importance for the sample.

### Tiny Illustrative Example

For a 2-feature sample  $x = (1.00, 2.00)$ ,  $y = 1$ ,  $\varepsilon = 10^{-3}$ . Suppose the RF gives:

$$p(x_1=1.001, x_2=2.00) = 0.82, \quad p(x_1=0.999, x_2=2.00) = 0.78.$$

Then  $\ell^+ = -\log(0.82) \approx 0.198$ ,  $\ell^- = -\log(0.78) \approx 0.248$ ,

$$g_1 \approx \frac{|0.198 - 0.248|}{0.002} \approx 25.0.$$

Similarly, if  $p(x_1=1.00, x_2=2.001) = 0.81$ ,  $p(x_1=1.00, x_2=1.999) = 0.805$ ,  $\ell^+ = 0.2107$ ,  $\ell^- = 0.2168$ ,

$$g_2 \approx \frac{|0.2107 - 0.2168|}{0.002} \approx 3.05.$$

Normalize:  $(\tilde{g}_1, \tilde{g}_2) \approx (0.891, 0.109)$ .

## 12 Reliability Criteria and Metrics

### 12.1 Gradient–Loss Reliability (GLR)

Let  $\mathbf{s}_i \in \mathbb{R}^d$  be the vector of *absolute* SHAP values for sample  $i$ , and let  $\tilde{\mathbf{g}}_i \in \mathbb{R}^d$  be the *normalized* finite-difference sensitivities (true loss sensitivity) for the same sample. We compute the per-sample **Spearman** rank correlation

$$\rho_i = \text{Spearman}(\mathbf{s}_i, \tilde{\mathbf{g}}_i),$$

and report the **GLR** score as the average across  $n$  samples,

$$\text{GLR} = \frac{1}{n} \sum_{i=1}^n \rho_i \in [-1, 1].$$

Higher is better: it indicates that SHAP’s ranking agrees with how the model’s loss locally reacts to feature nudges.

**What is Spearman correlation?** Spearman correlation (Spearman's  $\rho$ ) is a *rank correlation*: it measures agreement in *ordering*, not in numerical scale. This is ideal for feature importance because SHAP values and gradient/sensitivity magnitudes can live on different scales; if explanations are reliable, the *order* of important features should still match.

### Formula

Given vectors

$$u = (u_1, \dots, u_d), \quad v = (v_1, \dots, v_d),$$

first convert each to ranks:

$$r_j^u = \text{rank}(u_j), \quad r_j^v = \text{rank}(v_j).$$

Then compute the **Pearson correlation of the ranks**:

$$\rho = \frac{\sum_{j=1}^d (r_j^u - \bar{r}^u)(r_j^v - \bar{r}^v)}{\sqrt{\sum_{j=1}^d (r_j^u - \bar{r}^u)^2} \sqrt{\sum_{j=1}^d (r_j^v - \bar{r}^v)^2}}.$$

**Closed form without ties:**

$$\rho = 1 - \frac{6 \sum_{j=1}^d (r_j^u - r_j^v)^2}{d(d^2 - 1)}.$$

### Interpretation

- $\rho \approx 1$ : almost identical rankings  $\Rightarrow$  explanations consistent with loss sensitivity.
- $\rho \approx 0$ : no rank agreement  $\Rightarrow$  explanations do not reflect sensitivity.
- $\rho < 0$ : opposite rankings  $\Rightarrow$  explanations contradict sensitivity.

Thus, **GLR** summarizes explanation reliability as the mean Spearman  $\rho$  over samples.

## 12.2 $\varepsilon$ -Hit@k (Perturbation Validity)

**Definition (plain language).**  $\varepsilon$ -Hit@k tests whether the features that **SHAP** says are most important for a prediction are *actually* the ones that most increase the model's loss when we make a tiny change.

- “ $\varepsilon$ ” = a tiny feature nudge (e.g., +0.001).
- “Hit” = at least one of SHAP’s top- $k$  features also appears among the features causing the largest loss increases after nudging.
- “@ $k$ ” = how many top features we compare (top-1, top-3, top-5, …).

If the two top- $k$  sets overlap, the sample is a hit (1); otherwise 0. Averaging over samples yields the  $\varepsilon$ -Hit@k score.

**Step-by-step formulation.** For a single sample  $x \in \mathbb{R}^d$  with label  $y$ :

1. Let  $T_k$  be the indices of the **top- $k$**  features ranked by  $|\phi|$  (SHAP). In code we first restrict to the top  $M = \min(\text{m\_limit}, d)$  SHAP features and take  $T_k \subseteq \text{Top-}M$ .

2. For each candidate feature  $j \in \text{Top-}M$ , perturb one-at-a-time:

$$x^{+j} = x + \varepsilon e_j, \quad \Delta\ell_j = \ell(p(x^{+j}), y) - \ell(p(x), y),$$

where  $p(\cdot)$  is the model probability and  $\ell$  is BCE on probability.

3. Let  $D_k$  be the **top- $k$**  features ranked by  $\Delta\ell_j$ .

4. Define the per-sample indicator:

$$\text{Hit@k(sample)} = \mathbf{1}\{T_k \cap D_k \neq \emptyset\}.$$

Finally, over the evaluated samples,

$$\varepsilon\text{-Hit@k} = \frac{\#\text{hits}}{\#\text{tested samples}}.$$

### Pipeline (as implemented).

1. In each fold, randomly select  $\approx 15\text{--}20\%$  of test samples, respecting a **global budget** `max_eps_samples`.
2. For each chosen sample: compute SHAP  $\Rightarrow T_k$ ; for each of the top- $M$  SHAP features, add  $+\varepsilon$  and compute  $\Delta\ell_j \Rightarrow D_k$ ; record hit  $= \mathbf{1}\{T_k \cap D_k \neq \emptyset\}$ .
3. Average the hits across all chosen samples and folds to report  $\varepsilon\text{-Hit@k}$ .

**Simple numeric example.** One sample, three features  $(A, B, C)$ , true label  $y = 1$ ,  $k = 2$ ,  $\varepsilon = 0.001$ ,  $M \geq 3$ .

- **SHAP (absolute) rankings:**

$$|\phi_A| = 0.30, \quad |\phi_B| = 0.20, \quad |\phi_C| = 0.05 \Rightarrow \text{Top-}M = [A, B, C], \quad T_2 = \{A, B\}.$$

- **Loss after  $+\varepsilon$  perturbations:** (baseline loss = 0.20)

$$\begin{aligned} A : \ell(x + \varepsilon e_A) &= 0.30 \Rightarrow \Delta\ell_A = +0.10, \\ B : \ell(x + \varepsilon e_B) &= 0.25 \Rightarrow \Delta\ell_B = +0.05, \\ C : \ell(x + \varepsilon e_C) &= 0.205 \Rightarrow \Delta\ell_C = +0.005. \end{aligned}$$

Thus the  $\Delta\ell$  ranking is  $A > B > C \Rightarrow D_2 = \{A, B\}$ .

- **Compare:**  $T_2 = \{A, B\}$  and  $D_2 = \{A, B\}$  intersect  $\Rightarrow \text{Hit}@2(\text{sample}) = 1$ .

If instead  $D_2 = \{C, B\}$ , the intersection is  $\{B\} \Rightarrow$  still a hit; but if  $D_2 = \{C\}$ , then no overlap  $\Rightarrow$  hit = 0.

### 12.3 Calibration: Expected Calibration Error (ECE)

**What it measures.** The **Expected Calibration Error (ECE)** summarizes how well predicted probabilities match actual outcomes. A *well-calibrated* model that says “70% chance of defect” will be correct about 70% of the time for such predictions; an *over/under-confident* model will systematically over/under-shoot these frequencies.

**How it is computed.** Partition predicted probabilities into  $B$  uniform bins. In bin  $b$ , let  $\text{conf}(b)$  be the mean predicted probability and  $\text{acc}(b)$  the empirical positive rate. Then

$$\text{ECE} = \sum_{b=1}^B \frac{n_b}{N} |\text{acc}(b) - \text{conf}(b)|,$$

where  $N$  is the total number of samples and  $n_b$  is the number falling in bin  $b$ . Small  $\text{ECE} \approx 0$  indicates excellent calibration; larger values indicate worse calibration.

**Simple numeric example.** Consider 10 predictions and labels:

Sample	1	2	3	4	5	6	7	8	9	10
$\hat{p}$	0.10	0.15	0.20	0.30	0.40	0.55	0.60	0.70	0.80	0.90
$y$	0	0	0	1	0	1	1	1	1	1

Using  $B = 5$  bins:  $[0, 0.2)$ ,  $[0.2, 0.4)$ ,  $[0.4, 0.6)$ ,  $[0.6, 0.8)$ ,  $[0.8, 1.0]$ , we obtain:

$$\begin{aligned} \text{Bin 1: } n_1 = 2, \text{ conf}(1) = 0.125, \text{ acc}(1) = 0 &\Rightarrow \frac{n_1}{N} |\text{acc} - \text{conf}| = \frac{2}{10} \cdot |0 - 0.125| = 0.025, \\ \text{Bin 2: } n_2 = 3, \text{ conf}(2) = 0.300, \text{ acc}(2) = \frac{1}{3} \approx 0.333 &\Rightarrow \frac{3}{10} \cdot |0.333 - 0.300| = 0.0099, \\ \text{Bin 3: } n_3 = 1, \text{ conf}(3) = 0.55, \text{ acc}(3) = 1 &\Rightarrow \frac{1}{10} \cdot |1 - 0.55| = 0.045, \\ \text{Bin 4: } n_4 = 2, \text{ conf}(4) = 0.65, \text{ acc}(4) = 1 &\Rightarrow \frac{2}{10} \cdot |1 - 0.65| = 0.070, \\ \text{Bin 5: } n_5 = 2, \text{ conf}(5) = 0.85, \text{ acc}(5) = 1 &\Rightarrow \frac{2}{10} \cdot |1 - 0.85| = 0.030. \end{aligned}$$

Summing the contributions gives

$$\text{ECE} \approx 0.025 + 0.0099 + 0.045 + 0.070 + 0.030 = 0.18,$$

i.e., an 18% calibration error for this toy example. If  $\text{acc}(b) = \text{conf}(b)$  in every bin, then  $\text{ECE} = 0$  (perfect calibration).

## 12.4 Composite Reliability Score (reported in code)

We rescale GLR to  $[0, 1]$  via  $\rho' = (\rho + 1)/2$  and report

$$\text{ReliabilityScore} = \text{mean}\left(\rho', \varepsilon\text{-Hit@k}, \max\{0, 1 - \text{ECE}\}\right).$$

## 13 Experimental Setup

**Cross-validation.** We use **5-fold Stratified CV** with a fixed random seed (`rng=42`). Within each fold, SMOTE is applied *only to the training split* (adaptive  $k$ -neighbors) to address class imbalance, and a **Random Forest** (600 trees) is trained. Per-fold we compute test metrics (AUC, F1, Precision, Recall, Brier) and aggregate their means across folds.

**Artifacts (as produced by the code).** The pipeline returns:

- `per_fold_metrics`: table of fold-wise AUC/F1/Precision/Recall/Brier.
- `glr_rhos`: vector of per-sample Spearman  $\rho$  (SHAP vs. normalized loss sensitivity).
- `all_proba`, `all_y`: concatenated test probabilities and labels for calibration/ECE.
- `top_shap_mean_test`: mean absolute SHAP per feature (averaged across folds).

**Background for SHAP.** SHAP values on the test split are computed with `TreeExplainer` using an *interventional* background formed by a **random sample** of the training fold of size  $\min(256, |X_{\text{train}}|)$  (*not* K-means). Model output is set to **probability** and absolute SHAP values are used for importance.

**True-loss sensitivity & GLR.** For each test sample and feature, we estimate  $|\partial \ell / \partial x_j|$  via central finite differences on the BCE loss with a small perturbation  $\varepsilon = 10^{-3}$ ; sensitivities are row-normalized to  $\tilde{g}$ . **GLR** is the per-sample Spearman rank correlation between absolute SHAP and  $\tilde{g}$ , averaged over all test samples.

**$\varepsilon$ -Hit@k sampling.** To keep perturbation testing efficient, in each fold we select a  $\sim 15\text{--}20\%$  random subset of test samples, subject to a **global budget** `max_eps_samples` = 50. For each selected sample, we take the top- $M$  SHAP features ( $M = \min(\text{m\_limit}, d)$ ), nudge each by  $+\varepsilon$ , compute  $\Delta\ell$ , and record a hit if SHAP top- $k$  overlaps the  $\Delta\ell$  top- $k$ . The reported  **$\varepsilon$ -Hit@k** is the mean hit rate over sampled rows.

**Calibration.** Using all concatenated test probabilities/labels, we compute **ECE** with uniform bins (`calib_bins` = 10) and plot a reliability diagram via `calibration_curve`. A composite **ReliabilityScore** averages a rescaled GLR,  $\varepsilon$ -Hit@k, and  $(1 - \text{ECE})$  (clipped at 0) for a single summary measure.

## 14 Results

From the current run (summary row):

### Discrimination

- **AUC (mean across folds):** 0.7315
- **F1 (mean):** 0.2349
- **Precision (mean):** 0.2982
- **Recall (mean):** 0.2444
- **Brier (mean):** 0.1322

### Reliability of Explanations & Calibration

- **GLR (mean Spearman  $\rho$ ):** 0.0742
- **$\varepsilon$ -Hit@10:** 1.0
- **ECE:** 0.1098
- **Composite ReliabilityScore:** 0.8091
- **eps\_samples\_used:** 50 (global perturbation budget reached)

**Notes.** AUC indicates moderate ranking performance; F1/Recall are low, consistent with class imbalance. ECE  $\approx 0.11$  suggests reasonably calibrated probabilities. GLR is small ( $\sim 0.07$ ), implying weak overall rank alignment between SHAP and true-loss sensitivity on average, whereas Hit@10 is maximal (likely aided by a large  $k = 10$  relative to  $M$ ), yielding a high composite ReliabilityScore.